

EECE 5550 Mobile Robotics Lab #3

David M. Rosen

Due: Oct 31, 2022

Problem 1: Dynamic programming for Traveling Salesman Problem

The traveling salesman problem (TSP) is a problem in graph theory requiring to find the shortest possible route that visits each city exactly once and returns to the origin city. In this problem, we will consider an undirected graph containing four cities (A, B, C, D) as in Fig. 1. Here, an undirected graph implies that the weight on an edge is associated with both transitions (e.g., the costs of going from A to B and B to A are the same.)

Let A be the origin city. Then, all possible routes that start and end with A and visit all the other cities only once are: $ABCD A$, $ABDCA$, $ACBDA$, $ACDBA$, $ADBCA$, $ADCBA$.




Figure 1: A graph where A, B, C, D are the cities, the edges denote the transitions between the corresponding cities, and the weights on the edges refer to the cost of travel between the cities.

Note that the following tree can be generated for the route planning.

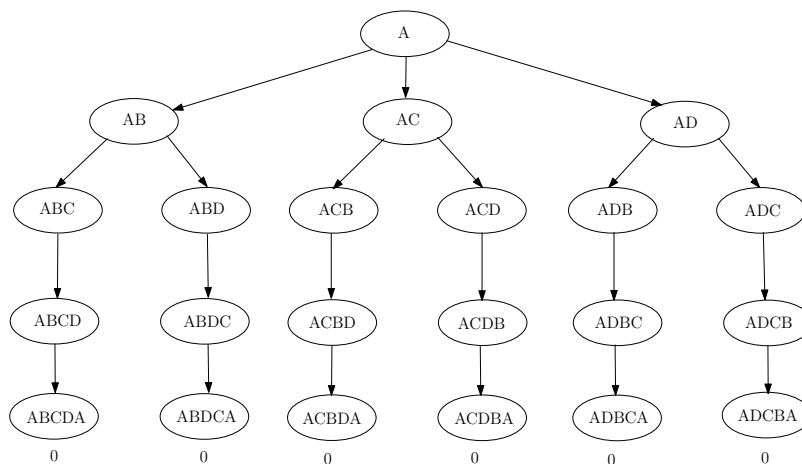



Figure 2: The decision-making tree where the cost-to-go of final states are zero.

- a. Use dynamic programming to compute the cost-to-go of each state on Fig. 2.
- b. By using your results from part a, for each possible route, identify an interval for α that makes the route optimal; or state if such an interval does not exist.


Problem 2: Route planning in occupancy grid maps

We saw in class that occupancy grid maps provide a convenient representation of a robot's environment that is particularly well-suited to route planning for navigation.

For example, Fig 3a shows a (probabilistic) occupancy grid map of a research lab constructed using the [Cartographer](#) SLAM system¹, and Fig. 3b the resulting estimate of free and occupied space obtained by thresholding these occupancy probabilities to binary values.



(a) Probabilistic occupancy grid map



(b) Binary occupancy grid map

In this exercise, you will implement two of the graph-based planning algorithms that we discussed in class (A* search and probabilistic road maps) to perform route-planning in the (binary) occupancy grid map shown in Fig. 3b.

Note: Since the A* search algorithm requires the use of several data structures other than basic matrices (e.g. sets and priority queues), we recommend implementing the following exercise in a Python notebook.

- (a) **A* search:** In the first part of this exercise, you will implement a general version of A* search that is abstracted with respect to the choice of representation of the graph G . This enables you to apply it to *both* occupancy grids (considered as 8-connected graphs) *and* “standard” graphs G (for use with probabilistic roadmaps).

The pseudocode for this version of A* search is shown as Algorithm 1. The following objects appear in Algorithm 1:

- “CostTo” is a [map](#) that assigns to each vertex v the cost of the shortest known path from the start node s to v .

¹We will see later in the course how to construct these maps from raw sensor data.

- “pred” is a map that associates to each vertex v its predecessor on the shortest known path from the start s to v .
- “EstTotalCost” is a map that assigns to each vertex v the sum $\text{CostTo}(v) + h(v, g)$, the sum of the cost of the best known path to v and the predicted cost of the best path from v to the goal g ; this is the estimated cost of the optimal path from the start s to the goal g that passes through vertex v .
- Q is a priority queue in which elements with *lower* values are removed *first*.
- “RecoverPath” is a function that takes as input the start state s , goal state g , and populated predecessor map pred, and returns the sequence of vertices on the optimal path from s to g .

Tip: In Python, you can use the [dictionary](#) class to implement the maps CostTo, pred, and EstTotalCost, a [list](#) or [set](#) to hold the vertex set V , and a sorted list or [heap](#) of (priority, vertex) tuples² to implement the priority queue Q .

- (i) Implement the RecoverPath function.
 - (ii) Using your implementation of RecoverPath, implement the complete A* search algorithm shown in Algorithm 1. Your code should accept as input the vertex set V , the start and goal vertices s and g , and function handles for N , w , and h .
- (b) **Route planning in occupancy grids with A* search:** In this part of the problem, you will apply your A* search algorithm to perform route planning directly on the occupancy grid map Fig. 3b, considered as an 8-connected graph. Note that in this part of the exercise, we will *identify* each vertex $v \in V$ (cell) using its row and column (r, c) in the occupancy grid.
- (i) Given an occupancy grid map M in the form of a 2D binary array, where a value of 0 indicates “occupied” and a value of 1 indicates “free” space, implement the function $N(v)$ that returns the set of unoccupied neighbors of a vertex v (remember that we can’t drive the robot through occupied space!) The vertex v and its neighbors should be expressed in the form of (row, column) tuples $v = (r, c)$.
 - (ii) We will consider the cost of moving from a cell $v_1 = (r_1, c_1)$ to an adjacent cell $v_2 = (r_2, c_2)$ to be the Euclidean distance between the cell centers. Implement a function $d: V \times V \rightarrow \mathbb{R}_+$ that accepts as input the tuples v_1 and v_2 , and returns this Euclidean distance.
 - (iii) We saw in class that the straight-line Euclidean distance between two points provides an admissible A* heuristic h for route-planning using the total path length as the cost; this means that we can do route planning using your distance function d from part (ii) as both the edge weight w and the heuristic h .

Using your implementations of d , N , and A* search, find the shortest path in the occupancy grid Fig. 3b from the starting point $s = (635, 140)$ to the goal $g = (350, 400)$ [assuming 0-based indexing for rows and columns, as is standard in CS.] Plot this optimal path overlaid on the image, and calculate its total length.

Tip: In Python, you can use the [Python Imaging Library](#) to easily manipulate basic image data. The following code snippet will read the occupancy map file from disk,

²Python compares tuples in lexicographic order, so placing the priority first ensures that (priority, vertex) tuples will be sorted by priority, as desired.

Algorithm 1 An abstracted implementation of A^* search

Input: Graph $G = (V, E)$ with vertex set V and edge set E , nonnegative weight function $w: V \times V \rightarrow \mathbb{R}_+$ for the edges, admissible A^* heuristic $h: V \times V \rightarrow \mathbb{R}_+$ that returns the estimated cost-to-go, a set-valued function $N: V \rightarrow 2^V$ that returns the neighbors of a vertex v in G , starting vertex $s \in V$, goal vertex $g \in V$.

Output: A least-cost path from s to g if one exists, or the empty set \emptyset if no path exists.

```
1: function A_STAR_SEARCH( $V, s, g, N, w, h$ )
    // Initialization
2:   for  $v \in V$  do
3:     Set  $\text{CostTo}[v] = +\infty$ .
4:     Set  $\text{EstTotalCost}[v] = +\infty$ .
5:   end for
6:   Set  $\text{CostTo}[s] = 0$ .                                 $\triangleright$  Cost to reach starting vertex  $s$  is 0.
7:   Set  $\text{EstTotalCost}[s] = h(s, g)$ .                     $\triangleright$  Estimated cost-to-go from  $s$  to  $g$ 
8:   Initialize  $Q = \{(s, h(s, g))\}$ .                   $\triangleright$  Insert start vertex  $s$  with value  $h(s, g)$ 
    // Main loop
9:   while  $Q$  is not empty do
10:     $v = Q.\text{pop}()$                                  $\triangleright$  Remove least-value element from  $Q$ 
11:    if  $v = g$  then                                 $\triangleright$  We have reached the goal!
12:      return RECOVERPATH( $s, g, \text{pred}$ )               $\triangleright$  Reconstruct and return optimal path
13:    end if
14:    for  $i \in N(v)$  do                                 $\triangleright$  For each of  $v$ 's neighbors
15:       $\text{pvi} = \text{CostTo}[v] + w(v, i)$                    $\triangleright$  Cost of path to reach  $i$  through  $v$ 
16:      if  $\text{pvi} < \text{CostTo}[i]$  then
17:        // The path to  $i$  through  $v$  is better than the previously-known best path to  $i$ ,
18:        // so record it as the new best path to  $i$ .
19:        Update  $\text{pred}[i] = v$ 
20:        Update  $\text{CostTo}[i] = \text{pvi}$                      $\triangleright$  Update cost of best path to  $i$ 
21:        Update  $\text{EstTotalCost}[i] = \text{pvi} + h(i, g)$ 
22:        if  $Q$  contains  $i$  then
23:           $Q.\text{setPriority}(i) = \text{EstTotalCost}[i]$        $\triangleright$  Update  $i$ 's priority
24:        else
25:           $Q.\text{insert}(i, \text{EstTotalCost}[i])$            $\triangleright$  Insert  $i$  into  $Q$  with priority  $\text{EstTotalCost}[i]$ 
26:        end if
27:      end for
28:    end while
29:  return  $\emptyset$                                  $\triangleright$  Return empty set: there is no path to goal
30: end function
```

interpret it as a Python `numpy` array, and then threshold it to produce the binary array M required in part (i):

```
# Load the PIL and numpy libraries
from PIL import Image
import numpy as np

# Read image from disk using PIL
occupancy_map_img = Image.open('occupancy_map.png')
```

```
# Interpret this image as a numpy array, and threshold its values to
↪ {0,1}
occupancy_grid = (np.asarray(occupancy_map_img) > 0).astype(int)
```

- (c) **Route planning with probabilistic roadmaps:** Voxelized grids (like occupancy maps) provide simple and convenient models of robot configuration spaces, but as we saw in class their memory requirements scale *exponentially* in the dimension of the state space, making them far too costly to use for higher-dimensional planning problems.

As we saw in class, *sampling-based planners* provide a tractable alternative for planning in high-dimensional spaces. Recall that these methods *approximate* the configuration space C using a graph $G = (V, E)$ whose vertex set $V \subset C$ is a *randomly sampled subset* of points in C , and where two vertices $v_1, v_2 \in V$ are joined by an edge if v_2 is *reachable* from v_1 by applying a local controller.

In this part of the exercise, you will implement a sampling-based planner to perform route planning in the occupancy grid shown in Fig. 3b; more specifically, you will implement a *probabilistic roadmap* (PRM). Recall that we construct a PRM incrementally by sampling a new vertex $v_{new} \in C$, and then attempting to join v_{new} to nearby vertices $v \in G$ using a local planner (cf. Algorithms 2 and 3). In order to implement this approach, we must therefore specify:

- A method for sampling new vertices $v_{new} \in C$ (line 4 of Alg. 2)
- A suitable distance function $d: V \times V \rightarrow \mathbb{R}_+$ for characterizing “nearby” vertices (line 3 of Alg. 3)
- A local planner (line 4 of Alg. 3)

Algorithm 2 Construction of a probabilistic roadmap

Input: Desired number of sample points N , maximum local search radius d_{max} .

Output: A graph $G = (V, E)$ consisting of a vertex set $V \subseteq C$ of cardinality N , and edge set E indicating reachability via local control.

```
1: function CONSTRUCTPRM( $N, d_{max}$ )
2:   Initialize  $V = \emptyset, E = \emptyset$ .
3:   for  $k = 1, \dots, N$  do
4:     Sample a new vertex  $v_{new} \in C$ .
5:     ADDVERTEX( $G, v_{new}, d_{max}$ )
6:   end for
7:   return  $G = (V, E)$ 
8: end function
```

- (i) Implement a function that accepts as input the occupancy grid map M , and returns a vertex $v = (r, c)$ sampled *uniformly randomly* from the free space in M . [Hint: consider rejection sampling with a uniform proposal distribution.]
- (ii) We saw in class that is easy to plan straight-line paths between arbitrary points using a differential drive robot, since the robot can rotate in-place to face the correct direction before beginning to move. Therefore, we might consider using a *straight-line path planner* as our local planner in line 4 of Alg. 3. Using this approach, a point $v_2 \in V$ is

Algorithm 3 Adding a vertex v_{new} to the probabilistic roadmap $G = (V, E)$

```
1: function ADDVERTEX( $G, v_{new}, d_{max}$ )
2:    $V \leftarrow V \cup \{v_{new}\}$ . ▷ Add vertex  $v_{new}$  to  $G$ 
3:   for  $v \in V$  satisfying  $v \neq v_{new}$  and  $d(v, v_{new}) \leq d_{max}$  do ▷ Link  $v_{new}$  to nearby vertices
4:     Attempt to plan a path from  $v_{new}$  to  $v$ .
5:     if planning succeeds then
6:        $E \leftarrow E \cup \{(v, v_{new})\}$  ▷ Add edge  $e = (v, v_{new})$  to  $G$ 
7:     end if
8:   end for
9: end function
```

reachable from v_1 if and only if the line segment joining v_1 and v_2 does not intersect any occupied cells in M .

Implement a function that performs this reachability check. Your function should accept as input the occupancy grid map M and two grid cells $v_1 = (r_1, c_1)$ and $v_2 = (r_2, c_2)$, and return a Boolean value indicating whether the line segment joining v_1 and v_2 in M is obstacle-free.

- (iii) With the aid of your results from parts (c)(i) and (c)(ii), and your distance function implementation from (b)(iii), implement Algorithm 2 in the form of a function that accepts as input an occupancy grid map M , the desired number of samples N , and the maximum local search radius d_{max} , and returns a PRM G constructed from M .

Tip: You may find it convenient to model the PRM G using the [Graph](#) class in Python's [NetworkX](#) library. If you do so, you can record the location (row and column) of each vertex v by setting its `pos` attribute. Similarly, given any straight-line path joining two vertices $v_1, v_2 \in V$ found in line 4 of Alg. 3, you can store the length of this path as the *weight* of the edge $e_{12} = (v_1, v_2)$ joining v_1 and v_2 in G . The following code snippet provides a minimal working example:

```
# Import the NetworkX library
import networkx as nx

# Create empty graph
G = nx.Graph()

# Add vertex v1 at position (r1, c1)
G.add_node(1, {'pos' : (r1, c1)})

# Add vertex v2 at position (r2, c2)
G.add_node(2, {'pos' : (r2, c2)})

# Add an edge between v1 and v2 with weight w12
G.add_edge(1, 2, weight=w12)
```

- (iv) Using your implementation of Algorithm 2, construct a PRM on the occupancy grid in Fig. 3b with $N = 2500$ samples and a maximum local search radius of $d_{max} = 75$ voxels. Plot the resulting graph overlaid on Fig. 3b. [Hint: you may find NetworkX's [draw_networkx](#) function useful here.]
- (v) Recall that given a PRM G for a configuration space C , and start $s \in C$ and goal $g \in C$, we can plan a route from s to g by first *adding* s and g to the PRM, and then

searching for a shortest path from s to g in G .

Using the PRM you constructed in part (v), find a path from $s = (635, 140)$ to $g = (350, 400)$. [Note: If s and g initially lie in separate connected components of G , you may need to sample and add more vertices to G until s and g are path-connected.] Plot this path overlaid on Fig. 3b, and calculate its total length.

Tip: In part (v), you may use NetworkX's implementation of [A* search](#).